



Snapdragon 888 / Qwen2.5-0.5B HMX FP16 — 12 outputs per independent sequence

Control Optimized


Aggregate backend decode throughput




HMX activity: positive returned samples



Peak SDK HMX active/load ratio



Peak SDK HVX active/load ratio



Sequential observed cases and repeats; no confidence intervals or statistical speedup claim.
 Decode excludes each first prefill-derived token. Counters span the whole capture, including initialization/prefill.
 SDK active/load percentages do not establish wall-time engine utilization or matrix row occupancy.
 Separate distinct-prompt equality check: 29/32 exact streams; numerical limitation retained in JSON.